

Bhavin Ingle

inglebhavin98@icloud.com • +917567288994 • linkedin.com/in/inglebhavin98

EDUCATION

Faculty of Management Studies, University of Delhi

Master of Business Administration (MBA)

Delhi

July 2024 – May 2026

Sardar Patel Institute of Technology, University of Mumbai

Bachelor of Engineering (BE) in Electronics & Telecommunication Engineering

Mumbai

June 2016 – May 2020

- Awarded Departmental and Institute Excellence honors for performance in Mitsubishi Electric Cup 2019; Led a team of 4 to rank 5/35 national finalists, secured INR 2L+ in corporate sponsorship and advanced PLC/HMI training.
- Secured INR 5L grant from Department of Science & Technology (GOI), Texas Instruments, NSRCEL (IIM-B), and filed 3 prov. patents for developing an aquaculture system; Top 30 Finalist team nationally among 100k+ applicants.
- Authored and presented 3 IEEE research papers (at IEEE Region 10 Conference – TENCON 2019, iSES, and GCAT) covering topics like cloud-based OPC telemetry, Industrial IoT automation architectures and computer vision.

EXPERIENCE

Management Consulting Analyst Intern

Accenture Strategy & Consulting

Bengaluru, IN

Apr 2024 – May 2024

- Architected a stateful multi-agent orchestrator prototype using LangGraph and LangChain to simulate the end-to-end migration of a Fortune 10 Global Healthcare Enterprise.
- Implemented an Orchestrator-Worker pattern running Chain-of-Thought (CoT) and few-shot prompting over multi-LLM router pipelines, mapping ~70 complex healthcare intents onto intelligent conversational steering bots.
- Engineered an automated data extraction and ingestion pipeline during a pre-signing client data freeze, scraping 100+ public enterprise URLs and unstructured policy PDFs to bypass operational project blockers.
- Deployed semantic and paragraph-based chunking protocols using LangChain to populate a low-latency ChromaDB vector database, providing a robust, ground-truth knowledge layer for context-aware RAG retrieval.

Senior Digital Engineer I

Tata Neu

Mumbai, IN

Sept 2020 – June 2024

- Led a 3-engineer pod to architect Cart and Checkout journeys with a hybrid Croma payment integration, driving a ~47% landing-to-checkout conversion increase, processing 14K+ weekly orders, and support 120Mn+ end users.
- Spearheaded Adobe Experience Manager (AEM) and React integration, decoupling frontend rendering from backend logic, cutting content delivery from days to minutes and empowering product teams to autonomously run A/B tests.
- Engineered aggressive WebView preloading and pre-rendering architectures within the Flutter superapp, eliminating browser engine cold starts and reducing Time-to-Interactive (TTI) by up to 70%.
- Built the 'My Orders' frontend module as sole owner, scaling it to top 5 most visited page during the 800K+ user rollout; earned 'The Spark' Award and 'Exceptional Achiever' rating for ranking in top 3 out of 250+ contributors.
- Fast-tracked for the Tata Administrative Services (TAS) flagship leadership program; shortlisted for the final Chairman's Group Discussion among 80+ top-performing internal candidates.

FOUNDING EXPERIENCE

Co-Founder

Zargun Electronics and Science Pvt. Ltd.

Remote

Jan 2018 – Aug 2020

Directed product engineering until MVP launch before transitioning to a non-operational advisory role in August 2020.

- Architected modular and ultra-low-powered IoT ecosystem devices with interchangeable sensor payloads, micro-spectrometry (to identify polymers – HDPE, PET, PP) enabling scaling across waste, compost, and material sensing.
- Recognized globally as a Top 15 enterprise (out of 41 teams across 16 nations) at the Singapore International Foundation's Young Social Entrepreneurs cohort, receiving mentorship from McKinsey and Temasek International.
- Secured over \$44K in seed grants, via the MeitY STPI CHUNAUTI challenge, alongside strategic victories at the IITB-SJMSOM Pitch 2.0, NXP India Tech Startup Challenge (Top 10), and Maharashtra Startup Week (Top 100).

PROJECTS

FlashInfer

Open-source library for high-performance LLM inference kernels

Jan 2023 – present

- Achieved 2.8x speedup over baseline attention implementations on A100 GPUs
- Adopted by 3 major AI labs, 8,500+ GitHub stars, 200+ contributors

SKILLS

Languages: Python, C++, CUDA, Rust, Julia

ML Frameworks: PyTorch, JAX, TensorFlow, Triton, ONNX

Infrastructure: Kubernetes, Ray, distributed training, AWS, GCP

Research Areas: Neural architecture search, model compression, efficient inference, multi-agent RL